

Biweekly Update — STF AI Diagnosis Platform

Period: 26 May – 9 June 2026

Note: A light period by design — I was heavily occupied with work and personal commitments, so **no milestone-level progress** was planned or made. Activity was **maintenance and collaborator-driven**: keeping the platform healthy and unblocking the team while real vehicle data started flowing in.

What happened (incremental, not milestones)

1. Agent diagnosis now streams live

https://github.com/qmohsu/stf_ai_diagnosis_platform_v1/issues/123

The "Agent AI" tab no longer shows a frozen spinner — reasoning, investigation steps, and the final diagnosis now stream in real time. A UX fix, shipped and live.

2. First real OBD log from Deng Xiao — parsed; timestamp bug caught & fixed

https://github.com/qmohsu/stf_ai_diagnosis_platform_v1/issues/119

Deng Xiao's first real Yamaha road-test log surfaced a data-quality bug (a corrupt end-of-file record made a 27-minute trip read as ~12 days). Fixed so corrupt rows are safely skipped. Clean session now ready for the golden set.

3. 15 OBD goldens reviewed by Deng Xiao

https://github.com/qmohsu/stf_ai_diagnosis_platform_v1/issues/120

Deng Xiao completed expert scoring and comments on all 15 OBD-agent golden answers; more teammate reviews to follow. This is the input the evaluation work has been waiting on.

Next: resuming real progress — 3 priorities

With the data and expert-reviewed goldens now in place, the focus shifts to the evaluation & tuning track:

1. OBD agent eval + tune — HARNESS-25

https://github.com/qmohsu/stf_ai_diagnosis_platform_v1/issues/117

Score the OBD agent against Deng Xiao's reviewed Yamaha goldens; iterate to improve.

2. Full pipeline tuning — HARNESS-24

https://github.com/qmohsu/stf_ai_diagnosis_platform_v1/issues/116

Tune prompt + eval + regression across the end-to-end agent.

3. First baseline eval — HARNESS-23

https://github.com/qmohsu/stf_ai_diagnosis_platform_v1/issues/107

Establish the first quality baseline: manual agent vs RAG on 30 locked goldens.

Suggested sequence: #107 first — the baseline (manual agent vs RAG) is the yardstick everything else tunes against, so it precedes the tuning iterations in #117 / #116.